

Machine Learning - Measuring and Predicting Right-Wing Populism in Britain Since the
2008 Financial Crisis

BSc. Computer Science

School of Informatics and Engineering, University of Sussex

Project Supervisor: Dr Maxine Sherman

1st May 2026

Statement of originality

This report is submitted as part requirement for the degree of BSc Computer Science at the University of Sussex. It is the product of my own labour except where indicated in the text. The report may be freely copied and distributed provided the source is acknowledged. I hereby give permission for a copy of this report to be loaned out to students in future years.

Acknowledgements

Thanks to Dr Maxine Sherman for supervising me and helping me throughout this project

Table of Contents

List of Abbreviations	6
1 Introduction	6
1.1 Overview	6
1.2 Motivation	6
1.3 Aim and Objectives	7
Objectives.....	7
2 Professional and Ethical considerations	7
3 Literature Review	8
3.1 Defining Right-Wing Populism	8
3.2 Measuring Populist Attitudes.....	9
3.3 Demographic and Social Predictors of Right-Wing Populism	10
3.4 Right-Wing Populism in Britain and Cross-Time Analysis.....	11
3.5 Research Gap and Dissertation Contribution	12
4 Methods	12
4.1 Research Design and Data	12
4.2 Selection and Harmonisation of Attitude Items	13
4.3 Measurement Strategy and Component Construction	13
4.4 Score Construction and Label Definition.....	16
4.5 Harmonisation of Predictors	16
4.6 Model Frame and Predictor Specification.....	17
4.7 Predictive Modelling Strategy	18
4.8 Hyperparameter Tuning and Temporal Validation	19
4.9 Trend Analysis	20
4.10 Use of Survey Weights and Interpretation.....	20
5 Results	20
5.1 Measurement Reliability Results	21
5.2 Predictive Modelling Results	21
5.3 Trend Results.....	23
5.3.1 Yearly Prevalence	23
5.3.2 Main Predictors Over Time	24
5.3.3 Subgroup Differences.....	25
6 Discussion	26
6.1 Main Findings	26
6.2 What the Measure Captures	26

6.3 Understanding the Main Social Differences.....	27
6.4 Understanding Change Over Time.....	28
6.5 What the Model Performance Means	28
6.6 Main Limitations	29
6.7 Future Work.....	29
6.8 Final Remarks.....	30
<i>Appendix</i>	30
Bibliography.....	30
Original Project Proposal.....	33
Signed Ethical Compliance Sheet	34

List of Abbreviations

BSA: British Social Attitudes survey
RWP: Right-wing populism
NLP: Natural language processing
PCA: Principal components analysis
RMSE: Root mean squared error
MAE: Mean absolute error
BCS: British Computer Society

1 Introduction

1.1 Overview

Right-wing populism has grown in prominence in Britain since the 2008 financial crisis, a period marked by major political and economic shocks, including major austerity, Brexit, COVID-19, and the cost-of-living crisis. Right-wing populism is generally characterised by anti-elite sentiment, low trust in politicians, and a strong nationalist outlook often shaped by anxieties over immigration, globalisation, and social change. In this dissertation, I use machine learning to examine whether an individual's demographic and social background can predict whether they are likely to hold right-wing populist attitudes, and to track how the importance of these predictors has changed over time.

Using data from the British Social Attitudes survey across five waves between 2011 and 2024, I construct an attitudinal measure of right-wing populism from repeated survey items and test how well a trained machine learning model can predict it from demographic and social variables. I also examine whether the strongest predictors remain stable across time or change as Britain moves through different political and economic periods. This allows me to assess both predictive performance and change over time.

1.2 Motivation

This project is motivated by my longstanding interest in politics and by having grown up during this period of ever-increasing social and political polarisation. Over time, I have seen right-wing populism become a more visible and influential force within our politics. I am interested in understanding why this is happening, not only theoretically but also empirically. This dissertation therefore uses data science and machine-learning methods to examine which characteristics are most strongly associated with

right-wing populist attitudes, how these attitudes are distributed, and how they change over time.

The period after 2008 is especially important to this project because it marks the start of a period of substantial political and economic change in Britain. Although the first wave I use is 2011, the post-2008 period provides a useful time frame because the financial crisis had a major impact on British politics, trust in the government, and the wider conditions that helped shape the rise and visibility of RWP in the years that followed.

1.3 Aim and Objectives

The aim of this dissertation is to examine whether the BSA data set can be used to predict right wing populism, and to analyse how the main predictors change over time.

This dissertation addresses two research questions:

1. Can a harmonised set of demographic and social variables be used to predict RWP in the BSA?
2. Do the strongest predictors of RWP remain stable across survey waves, or do they change over time?

Objectives

- To construct a consistent measure of right-wing populism using harmonised attitudinal items from the BSA.
- To build a harmonised set of demographic and social predictors that can be used consistently across survey waves.
- To evaluate how well this predictor set predicts the final right-wing populism measure using ridge regression.
- To test whether predictive performance remains useful in a holdout wave
- To identify the strongest predictors of the measure and check whether their associations remain stable or change across the years.

2 Professional and Ethical considerations

This project was conducted in line with the professional principles set out by the BCS Code of Conduct, particularly those relating to public interest, privacy, competence, integrity, and accountability. The analysis was designed to minimise risk to individuals, avoid discriminatory or stigmatising interpretation, and report methods and findings transparently.

The project used secondary data from the British Social Attitudes (BSA) survey, collected by NatCen Social Research and accessed in anonymised form. No new data was collected from human participants, and no attempt was made to re-identify

respondents or link the survey data to external personal information. As a result, the project was low risk in ethical terms. The signed Ethical Compliance Form is included in the appendix.

Data was used in accordance with the relevant access conditions, University guidance, and applicable data protection principles. Because the project relied on politically and socially sensitive attitudinal data, particular care was taken in how results were framed and reported. The final right-wing populism measure used in the dissertation is an analytical construct derived from survey responses rather than a direct survey variable or a fixed statement about any individual. The 'RWP_top_20' label is a threshold-based summary derived from that measure, both are interpreted cautiously and discussed in neutral language.

The project also considered harm minimisation in its analytical design. Survey weights were used where appropriate to improve population-level inference, and limitations and uncertainty were reported clearly in order to avoid overstating the meaning or performance of the results. The analysis focused on broad trends across different survey waves rather than judgement about particular individuals or groups.

Finally, the project was carried out within the limits of the researcher's competence and with an emphasis on transparency and accountability. Analytical decisions, configuration choices, preprocessing steps, and modelling procedures were documented so that the code and project structure could be reviewed and reproduced. This was intended to support both methodological integrity and responsible reporting of findings.

3 Literature Review

3.1 Defining Right-Wing Populism

A useful starting point in the literature is Mudde's definition of populism as a thin-centred ideology that separates society into "the pure people" and "the corrupt elite", and holds that politics should express the general will of the people [1]. This definition has been influential because it gives populism a clear conceptual core while also explaining why it can attach itself to broader ideological traditions rather than operating as a fully developed ideology on its own [1].

This is important because right-wing populism is not simply populism in a general form. In the European literature it is more closely related to the populist radical right, which Mudde defines through the combination of populism, nativism, and authoritarianism [3]. On that view, the people vs elite logic remains central, but it is joined to a more exclusionary understanding of national community, social order, and belonging [3]. Right-wing populism therefore involves not only anti-elite conflict, but also a specifically right-wing emphasis on borders, hierarchy, and cultural exclusion [3].

At the same time, there is still disagreement over what populism actually is. Some scholars treat it less as a set of ideas and more as a political style, marked by performances of crisis, appeals to authenticity, and a claimed closeness to ordinary people [2]. This perspective is useful for understanding how populists present themselves and mobilise support, even if it is less directly suited to measuring populist attitudes through survey data.

It is also important to distinguish right-wing populism from several similar concepts. It is not the same as authoritarianism, since support for order, discipline, or punishment does not by itself contain the people vs elite moral conflict associated with populism [1, 3]. It is not identical to the radical right more broadly, because radical-right politics includes nativist and authoritarian elements that are not always expressed in explicitly populist terms [3]. Nor is right-wing populism reducible to just political distrust alone. Although distrust in politicians and institutions is often associated with populist attitudes, research shows that populist attitudes can still be separated empirically from political trust and from external political efficacy [7].

For this dissertation, right-wing populism is treated as an approximation built from the BSA items that are repeatedly available across waves. In practice, this means using a combination of distrust, authoritarian, immigration, and welfare-related attitudes to capture part of the broader idea discussed in the literature, while recognising that these items do not form a complete or fully valid populism scale. The measure should therefore be understood as more of an approximation constrained by item availability, rather than as a full reproduction of the wider theoretical concept.

3.2 Measuring Populist Attitudes

The measurement literature broadly agrees that populism can be studied through survey attitudes, but there is much less agreement about exactly how it should be measured. One important strand of research develops scales based on anti-elitism, people-centrism, and popular sovereignty. Akkerman, Mudde and Zaslove's scale was especially important because it showed that populist attitudes could be measured among ordinary voters, not only in parties, leaders, or manifestos [4]. Later work refined this further by arguing that populist attitudes should be treated as multidimensional rather than one-dimensional. Schulz and colleagues, for example, proposed a three-dimensional approach based on anti-elitism, support for popular sovereignty, and belief in the homogeneity and virtue of the people [5].

However, agreeing on dimensions does not resolve the main measurement problem. A central debate is whether populist attitudes should be treated as compensatory or non-compensatory. A compensatory approach would allow a high score on one component to offset a lower score on another. For example, a respondent with strong anti-elite attitudes might still be treated as highly populist even if they score less strongly on people-centrism. Wuttke, Schimpf and Schoen argue that this is conceptually weak because populism is not simply a loose bundle of traits. Instead, they describe it as an attitudinal syndrome in which the core elements are jointly necessary [6]. So, in a

sense, someone who is strongly anti-elite but not committed to popular sovereignty should not be treated as equally populist as someone who follows all of the core dimensions together [6].

This debate matters directly for this dissertation. The BSA does not provide a single dedicated set of populism items that can be used unchanged across all the waves analysed here. Instead, it provides a smaller set of repeated items that can be used to approximate a populist component and combine it with a right-authoritarian component. The literature therefore matters because it highlights the conceptual problems involved in collapsing all of the items into one broad score, and supports the need for a more structured and transparent measurement approach [6]. It also shows that measurement is not just a technical issue. It involves theoretical decisions about what should count as populism in the first place.

Measurement is also complicated by the problem of comparability across contexts and over time. A scale developed in one setting will not always transfer neatly to another, and even within one country there has to be some confidence that repeated items continue to capture the same underlying construct. Research on cross-national measurement invariance suggests that populist-attitudes inventories can show a common underlying structure, but only if the function of items and model specification are handled carefully [8]. Other studies show that some scales perform less well when populists are already in government, because anti-elite items may become harder to interpret in those contexts [9]. This matters here because any cross-wave analysis depends heavily on how populism is operationalised. In this dissertation, the aim is not to produce a perfect universal measure of populism, but to construct a repeatable, transparent, and conceptually defensible approximation that can be applied consistently across all the BSA waves [5, 6, 8, 9].

3.3 Demographic and Social Predictors of Right-Wing Populism

The literature on the social bases of right-wing populism is large, but some findings appear more consistently than others. Education is one of the strongest and most widely reported. Lower education is often associated with stronger support for populist radical right parties, and this is usually interpreted through arguments about cultural backlash, declining status security, or weaker attachment to liberal-cosmopolitan values [10]. Some research also suggests that this education divide can become stronger in areas with a larger immigrant presence [17]. Age also matters, but not always in a simple or consistent way. In the Brexit case, Leave support was especially common among older voters [11]. Other research suggests that age can also shape how strongly educational and rural-urban divides appear in radical right support, rather than acting only as a simple predictor on its own [23]. These variables do not explain populism on their own, but they do repeatedly point to important social divisions.

Class and income also matter, although the evidence here is less consistent. Some work highlights the role of economic insecurity, especially among people who feel left behind by economic restructuring, globalisation, or relative status loss [10, 18]. At the

same time, research suggests that the effects of economic decline and immigration are not always straightforward, and may vary by context and model specification [22]. The Brexit literature is especially useful here because it shows that support for Leave was concentrated among older, less educated voters, and was also strongly tied to concerns about immigration and cultural change [11]. This suggests that demographic variables may contain useful information about who is more likely to support right-wing populist politics, even if they do not provide a full explanation on their own.

Gender, religion, and region are more context dependent. In some settings men are somewhat more likely than women to express right-populist attitudes, and research on radical right populist voting has often found women less likely to support such parties, though the size of that gap varies across cases [19]. Religion can matter too, but its effect varies depending on context: in much of Western Europe more religious voters are not more likely to support populist radical right parties, while in parts of East-Central Europe the relationship can run in the opposite direction [20]. Regional differences are also difficult to interpret because they may reflect political context, economic inequalities, public service availability, and the clustering of other social characteristics within places. Research on urban-rural divides shows that support for far-right or populist radical right parties is often higher outside urban areas, although these differences also overlap with other social and economic factors [21]. Recent work also suggests that the strength of this rural-urban divide can vary by age, and that it can also overlap with education [23].

An important limit in this literature is that demographic variables are recurring indicators rather than full explanations. Recent work shows that populist attitudes are also strongly connected to trust, political efficacy, perceived deprivation, mistrust of experts, conspiracy beliefs, and simple “common sense” views of politics [15, 16]. This means that demographics can help identify broad social differences linked to right-wing populism, but they do not identify most of the deeper psychological or attitudinal mechanisms behind it. For that reason, demographic variables are useful for showing who is more likely to hold these attitudes, even if they cannot fully explain why such attitudes develop [10, 15, 16].

3.4 Right-Wing Populism in Britain and Cross-Time Analysis

Britain is a useful case for studying right-wing populism because the period covered in this dissertation includes a series of major economic and political shocks, including the aftermath of the financial crisis, Brexit, the pandemic, and the cost-of-living crisis. These developments intensified debates around sovereignty, immigration, and political representation, especially in the context of Brexit [11, 24]. They also took place alongside lower trust and confidence in British politics during the 2019–24 Parliament [12]. Britain is useful here not just because the data is available, but because it is a setting in which distrust of elites, identity politics, and questions of political authority have become especially prominent.

The Brexit referendum is particularly important in this respect. Hobolt's analysis of the Brexit vote highlights the strong divide between more and less educated voters, between younger and older voters, and between those more and less concerned about immigration and multiculturalism [11]. Leave support is not the same thing as right-wing populism, but the referendum did expose a wider conflict around national identity, distrust of establishment politics, and hostility to elites that overlaps clearly with themes in the right-populist literature [11].

This makes Britain a useful setting for studying change over time. Recent work on the stability of populist attitudes suggests that populism is neither a fully fixed trait nor a purely short-term fluctuation, which reinforces the importance of treating stability and change as empirical questions rather than assumptions [14].

3.5 Research Gap and Dissertation Contribution

The current literature has already done three things well. First, it has developed a clear conceptual discussion of populism and clarified its relationship to similar ideas such as authoritarianism and the radical right [1, 3]. Second, there is substantial measurement literature available showing how populist attitudes can be studied through surveys, while also exposing problems of dimensionality, aggregation, and invariance [4, 5, 6, 8, 9]. Third, it has identified a recurring set of demographic and attitudinal correlates associated with right-wing populism [10, 15, 16].

Although there are some explicitly predictive studies [15], what the literature has done less often is test how far a standardised demographic predictor set can actually predict a measure of right-wing populism using survey data collected across multiple waves. Much of the existing research is explanatory or correlational rather than explicitly predictive. There are also fewer studies that examine whether the importance of the same predictors changes over time under one consistent measurement approach, rather than assuming that pooled results reflect a stable social structure [14, 15].

This dissertation addresses that gap by constructing a repeatable survey-based right-wing populism measure from BSA data, testing how far a standardised demographic predictor set can predict that measure, and examining how predictor importance changes across waves. Its contribution is therefore not to settle the broader conceptual debate on populism, but to connect that debate to an empirical design that takes both measurement trade-offs and temporal change seriously.

4 Methods

4.1 Research Design and Data

This dissertation uses secondary data from the British Social Attitudes survey (BSA), a long-running survey of adults in Britain. The BSA is useful for showing how attitudes are distributed across the population and how they change over time. However, because the survey does not follow the same people over time, it cannot show how the same people changed from one wave to the next.

My analysis uses five waves: 2011, 2013, 2019, 2023, and 2024. These waves were chosen after checking the BSA waves from 2008 onward for repeated item availability. They contained the most relevant repeated items for building the RWP measure, they also cover a relatively consistent time period from 2008 to present day which makes them useful to analyse over time. Across these waves, the raw pooled dataset contains 19,477 respondents. After cleaning, harmonising variables, constructing the final score, and filtering to respondents with a valid outcome and survey weight, the final number of respondents I had was 17,615. 5,559 of these are in the 2023 wave, which I keep as a dedicated holdout validation year. The remaining four waves are used to build the RWP measure, and train and test the model.

4.2 Selection and Harmonisation of Attitude Items

Harmonisation means making data measured in slightly different ways comparable enough to be used together in one analysis.

The attitudinal side of the project is built from repeated BSA questions chosen for the analysis. Before harmonising them, I first needed to decide which questions were worth using. I used NLP to help find candidate attitude questions from the survey documentation, because there were too many possible items to identify them all manually in a consistent way. This was only used to help narrow down the set of possible items. Final item choice was based on relevance, repeated availability across years, and the later measurement checks.

Each selected item was then put onto a common 0 to 1 scale, and reverse-coded where needed so that higher values always meant more of the attitude being measured. Invalid codes, missing codes, excluded codes, and out-of-range values were set to missing before rescaling. This does not make every item exactly the same across time, but it does make them easier to compare and score consistently. I also produced per-wave checks for each attitude item so that missingness and item coverage could be reviewed before fitting the measurement model.


4.3 Measurement Strategy and Component Construction

The BSA does not contain a ready-made measure of RWP, so the measure used in this dissertation is built from repeated items that are available across the selected waves, which is why it should be viewed as an approximation, rather than a perfect definition of right-wing populism.


I use principal components analysis (PCA) to help identify how these items group together. PCA is a statistical method used to reduce a larger set of related variables into a smaller number of broader dimensions. It does this by identifying which variables tend to move together.

Before the PCA I start from one combined pool of RWP-related items. This pool includes distrust, authoritarian, immigration, and welfare items. The items are then checked in two ways. First, an item has to be available in all four non-holdout reference years used: 2011, 2013, 2019, and 2024. Then I remove any item that is very weakly related to the other items in its starting group, using a strongest absolute correlation cutoff of 0.1. The remaining items are then pooled and entered into the PCA.

Not everyone answered every question, so I could not run PCA directly on the raw respondent item table. Instead, I first calculated how strongly the items were related to each other using the available responses and then ran PCA on that item-correlation matrix. Missing answers do not stop someone contributing to the PCA, because the PCA is built from the available item relationships, not from complete respondent rows only. The 2023 holdout year is not used in this step. Varimax rotation is then used to make the loadings easier to interpret.



The number of retained components is chosen using a scree-elbow rule. Items are only assigned to a component if they have an absolute loading of at least 0.3 and a loading gap of at least 0.1 over the next strongest component.



This produces two retained components. The first I label ‘right’ and combines distrust, authoritarian, and immigration items. The second I label ‘welfare’ and it contains welfare-related attitudes. The final retained structure contains 17 items, and those same items are available in all five study years: 2011, 2013, 2019, 2023, and 2024.

Retained Measurement Questions

Construct	Component	Question	Wave Years
Rwp	Right	Government trust	2011, 2013, 2019, 2024
Rwp	Right	MP trust	2011, 2013, 2019, 2024
Rwp	Right	Censorship morality	2011, 2013, 2019, 2024
Rwp	Right	Death penalty	2011, 2013, 2019, 2024
Rwp	Right	Immigration culture	2011, 2013, 2019, 2024
Rwp	Right	Immigration economy	2011, 2013, 2019, 2024
Rwp	Right	Obey authority	2011, 2013, 2019, 2024
Rwp	Right	Obey law	2011, 2013, 2019, 2024
Rwp	Right	Stiff sentences	2011, 2013, 2019, 2024
Rwp	Right	Traditional values	2011, 2013, 2019, 2024
Rwp	Welfare	Welfare damage lives	2011, 2013, 2019, 2024
Rwp	Welfare	Welfare deserving	2011, 2013, 2019, 2024
Rwp	Welfare	Welfare fiddling	2011, 2013, 2019, 2024
Rwp	Welfare	Welfare helping each other	2011, 2013, 2019, 2024
Rwp	Welfare	Welfare pride	2011, 2013, 2019, 2024
Rwp	Welfare	Welfare spending	2011, 2013, 2019, 2024
Rwp	Welfare	Welfare stand own feet	2011, 2013, 2019, 2024

I do not use PCA to create the final respondent scores directly. Instead, each respondent’s component scores are calculated by taking the mean of their answered 0 to 1 item responses within each retained component. I use this approach because it

keeps the final measure easier to interpret, while still using PCA to check the structure of the items.

As a final check, I also calculate complete-case Cronbach's alpha and an approximate omega total for each retained component in the pooled data and in each wave separately. These are standard reliability measures and are used to check whether the retained components are consistent enough to use in the later stages of the analysis.

4.4 Score Construction and Label Definition

For each respondent, I first calculate a 'right' score and a 'welfare' score by taking the mean of the answered items within each component. For example, if a respondent answered 8 'right' items and those values summed to 5.6, their 'right' score would be 5.6 divided by 8, which equals 0.70. I use this approach because all retained items have already been put onto the same 0 to 1 scale, so the final component scores stay on that same scale and remain easy to interpret.

Missing answers are not filled in at this stage. Instead, a component score is only calculated if the respondent answered at least 60% of the items in that component, rounded up to the nearest whole item. In the final retained structure, this means at least 6 of the 10 'right' items and at least 5 of the 7 'welfare' items. In practice, this means a respondent can miss a few questions and still receive a component score, but if they miss too many then that component is left as missing. I use this rule to avoid treating very partial responses as meaningful scores, while also avoiding the unnecessary loss of respondents who missed only a small number of items.

Once those component scores have been created, I calculate the final 0 to 1 outcome, 'rwp_score'. This is the average of the available retained component scores for each respondent. If both 'right' and 'welfare' are available, 'rwp_score' is the mean of the two. If only one of them is available, the final score is based on that one component alone, although this only affects about 1% of the final sample. If neither is available, 'rwp_score' is missing. I use this approach because the final measure is intended to summarise both retained parts of the construct, while still making use of the information available for each respondent rather than dropping cases unnecessarily.

I also create 'RWP_top_20' for the prevalence and trend outputs. It includes respondents with a 'rwp_score' above the weighted top 20% cutoff. I calculate this cutoff once, using the pooled non-holdout years: 2011, 2013, 2019, and 2024. I then apply the same cutoff to each year, so the threshold does not change between waves. The main modelling stage still uses 'rwp_score'.

4.5 Harmonisation of Predictors

To compare respondents across time, I harmonised a set of demographic predictors: age, year, gender, education, income, social class, region, religion, religiosity, ethnicity,

housing tenure, employment status, and marital status. Harmonisation here means making sure these variables could be used in a comparable way across the different survey years, even where the original categories were not exactly the same.

Age and year were kept as numeric variables rather than being grouped into categories. The remaining demographic variables were recoded into common categories. Education was grouped into low, mid, and high categories. The high group mainly includes degree-level and higher-education qualifications, the mid group includes A-levels or equivalent qualifications, and the low group includes qualifications below A-level and no qualifications. Responses where the qualification level could not be aligned clearly across years were treated as missing.

Employment status was grouped into employed, unemployed, retired, student, and inactive. Class was grouped into high, mid, low, and other so that it could capture broad social position consistently, while still leaving room for cases that did not fit neatly into the main three groups. Religion was grouped into no religion, Christian, and non-Christian. The other demographic variables were treated in the same general way so that they could be used consistently across waves. Where a value could not be aligned clearly across years, it was treated as missing rather than forced into a category that was difficult to justify.

I also checked each predictor year by year after harmonisation. I looked at high missingness or weaker comparability across years. Income stood out as the clearest example, as it had lower usable coverage than most other predictors in every wave. However, no predictor was missing enough to be removed from the final main predictor set.

4.6 Model Frame and Predictor Specification

The predictor and score data are joined into one modelling dataset so that each respondent's demographics are matched to their own 'rwp_score'. For each wave, I join the demographic predictor table to the score outcome table using respondent identifier and survey year. I checked that each join was one-to-one, so respondents were not duplicated or mismatched.

Before fitting the regression model, the modelling dataset is split into X and y. X contains only the demographic and social predictors, while y contains the 'rwp_score' scores that the model is trained to predict. The original attitude items and 'RWP_top_20' are not used as predictors in the main ridge model. I also tracked row loss from the raw wave size to the final sample, so I could see how many respondents were excluded because they did not have a valid 'rwp_score' or survey weight.

The final main predictor set contains 13 variables: age, year, gender, education, income, class, region, religion, religiosity, ethnicity, tenure, employment status, and marital status. I only kept predictors that were usable in at least three of the four waves. This allowed for one weaker wave, because missing predictor values are handled later during

preprocessing, but it removes any predictor that was only really available in one or two years. In the final run, all 13 predictors passed this check.

For the trend analysis, I used a stricter version of the same coverage check. A predictor only counted as usable in a wave if fewer than 40% of respondents were missing that predictor. I also excluded year from the trend predictor set, because the trend models are fitted separately within each year. This stricter rule was used because the trend analysis compares predictor importance across waves, so the predictors need stronger and more even coverage over time.

4.7 Predictive Modelling Strategy

The main predictive task in the project is to estimate where respondents sit on the 0 to 1 'rwp_score' using demographic and social predictors. In other words, the aim is to see how much information about a respondent's final score can be recovered from the predictor set alone.

Regression is a statistical method used to estimate the relationship between a set of predictors and a continuous outcome. In this project, the outcome is the 'rwp_score'. The final main model is ridge regression. Ridge regression is a form of linear regression that adds a penalty to large coefficients. The reason for using it here is that it helps keep the model more stable when predictors overlap with each other, when some are more informative than others, and when the model might otherwise put too much weight on small or noisy patterns in the data.

The main comparison is between a mean baseline and a ridge regression model. Both are fitted using survey weights. The baseline gives every respondent the weighted mean 'rwp_score' from the training data, while the ridge model uses the full demographic predictor set. This shows whether the predictors add useful information beyond a simple no-information benchmark.

The 2023 wave is removed before model development and kept aside as a final holdout test. The remaining rows from 2011, 2013, 2019, and 2024 are then split randomly into a 70/30 training-test split using a fixed seed of 42. In the final run, this produces 8,439 training rows and 3,617 internal test rows. All randomness used a fixed seed of 42 so that steps, such as the data split, could be reproduced consistently across runs.

In the modelling stage, the predictors are the input data that go into the model, while the outcome is the value the model is trying to predict. As I mentioned earlier a clear way to understand it is; the predictors are X and the outcome is y. Here, X contains the harmonised demographic and social predictors, while y is the 'rwp_score'.

I then apply a second missingness check to the predictor data in X, using the training data only. This is separate from the earlier wave-coverage check. The earlier check tests whether a predictor is available across enough years, while this second check tests whether it has too much missing data in the actual training set. If a predictor has more

than 70% missingness in the training data, it is removed before model fitting. In the final run, no predictors were removed by either rule. Also respondents only enter the modelling sample if they have a valid 'rwp_score' and survey weight.

Before fitting the models, the predictor data in X are preprocessed so that the regression can use them properly. Missing numeric predictor values are filled in using the mean of that predictor. although in practice numeric missingness was very low: in the training data, only age had missing values, affecting about 0.3% of rows. Missing categorical predictor values are just coded as 'missing', which mattered more because some categorical predictors had noticeably higher missingness.

The categorical predictors in X are then converted into 0/1 columns using one-hot encoding. For example, education is represented with columns for categories such as 'mid' and 'high', while 'low' is left out as the baseline reference group. This means the coefficients for 'mid' and 'high' are interpreted relative to 'low'.

Model performance is evaluated on both the internal test split and the 2023 holdout. I report RMSE, MAE, and R^2 in both weighted and unweighted form for both evaluations. For the internal test split, I also report Pearson and Spearman correlations, compare the ridge model against the mean baseline, and include a predicted standard deviation ratio to show how compressed the predictions are relative to the observed scores.

4.8 Hyperparameter Tuning and Temporal Validation

A hyperparameter is a setting chosen before the model is fitted, rather than something the model learns directly from the data. In ridge regression, the main hyperparameter is 'alpha'. This controls how strongly the model shrinks its coefficients. Lower values allow the model to fit the data more freely, while higher values shrink the coefficients more strongly. I tune this parameter because it affects how flexible or constrained the final model is. I tune 'alpha' using a small fixed grid of candidate values: 1, 3, 10, 30, 100, 300, 1000, and 3000. This tuning is carried out on the non-holdout data only, so the 2023 holdout plays no role in choosing the parameter.

I use the training split for five-fold cross-validation. This means the training data are split into five parts, and the model is repeatedly fit on four parts and checked on the fifth, so that each part is used for validation once. Wave year is used as the stratification target for this step. Candidate values are ranked mainly by mean weighted RMSE, because lower prediction error is better. Weighted MAE is then used as a tie-breaker, and if there is still a tie, the higher 'alpha' is preferred. The value selected by cross-validation is then carried forward into the final model, while the 2023 wave is kept aside for final temporal evaluation only. In the final version, the retained value is 'alpha = 30'.

Temporal validation is handled through the dedicated 2023 holdout. This wave is excluded from measurement fitting, threshold estimation, model fitting, and internal train-test evaluation. The holdout therefore acts as a final test of whether relationships learned from 2011, 2013, 2019, and 2024 still generalise to a different survey wave.

4.9 Trend Analysis

The project also includes a separate trend-analysis stage, which is descriptive rather than predictive. The 2023 holdout year is excluded from these outputs. I use 'RWP_top_20' to show yearly prevalence, meaning the share of respondents above the top-20 cutoff in each wave. I then fit separate linear regression models within each year using 'rwp_score'. This is done instead of reusing the pooled ridge model because the aim is to describe yearly associations, not to make predictions.

To compare predictor strength across years, I use grouped permutation importance. This works by shuffling one original predictor at a time while leaving the others unchanged, and then measuring how much the model's weighted R^2 falls as a result. A larger drop in weighted R^2 means the model relied more heavily on that predictor in that year. This method is useful because many categorical predictors are split into several encoded columns during preprocessing, so grouping them under the original predictor allows each variable to be assessed as a single predictor across years.

4.10 Use of Survey Weights and Interpretation

Survey weights are used throughout the project whenever the aim is to say something about the wider population, rather than just the sample that happened to be observed. In the BSA, these weights help adjust for the survey design and make the results more representative. In this project, that means weights are used in threshold estimation, binary label construction, prevalence estimation, ridge regression fitting, baseline comparisons, internal and holdout evaluation, year-specific trend regressions, and subgroup summaries. The main exception is the PCA stage. PCA is fitted on an unweighted item-correlation matrix and is used here to organise the structure of the attitude items, rather than to produce population estimates. At that stage, the aim is just to see which items group together, not to estimate how common those attitudes are in the wider population.

Because the BSA surveys different people in different years, this dissertation can show associations but not causality. It can show which predictors tend to go with higher or lower 'rwp_score', but not whether those predictors cause the attitudes. Since 'rwp_score' is constructed from the available BSA questions, and 'RWP_top_20' is only the top-20 group created from that score, both should be treated as practical tools for this project rather than exact measures of a person's true level of RWP.

5 Results

5.1 Measurement Reliability Results

Cronbach's alpha and omega are both measures of internal consistency. In other words, they show whether the items in each component fit together well enough to be used as part of the same scale. The results give mixed but usable evidence for the retained components. In the pooled data, Cronbach's alpha was about 0.80 for 'right' and 0.85 for 'welfare', while omega was lower at about 0.59 and 0.68. This means the components show reasonable consistency by alpha, but the omega results suggest that the measure should still be treated cautiously. Overall, the final measure looks good enough to use but is by no means a perfect representation of right-wing populism.

5.2 Predictive Modelling Results

The first question to answer was how well 'rwp_score' could be predicted from demographic and social predictors. To assess this, I compared the ridge regression model against a mean baseline. The mean baseline is the simplest possible comparison, because it gives every respondent the same predicted score. This means it acts as a minimum benchmark: if the ridge model cannot improve on it, then the predictors are not adding useful information.


prediction performance summary

Metric	Test split	Holdout
Model	ridge_regression	
Predictors	13	
Test rows	3617	
Holdout year		2023
Holdout rows		5559
Weighted R2	0.159788	0.143167
Weighted RMSE	0.145647	0.168646
Weighted MAE	0.11434	0.132682
Pearson	0.415257	
Spearman	0.397474	
Predicted SD ratio	0.375001	
RMSE gain vs mean baseline	8.3%	
MAE gain vs mean baseline	7.6%	

The ridge model did improve on that benchmark. On the internal test split, it achieved a weighted R^2 of 0.16. This means it explained around 16% of the variation in 'rwp_score'. On the 2023 holdout, the weighted R^2 was 0.143, so performance was slightly weaker on unseen data, but not by much. This suggests that the model is finding patterns that carry over to a different wave, rather than only fitting the years it was trained on.

The error measures paint the same general picture. On the internal test split, weighted RMSE was 0.146 and weighted MAE was 0.114. On the holdout, these rose to 0.169 and

0.133. So the model is less accurate on the holdout than on the internal test data, but not dramatically so. Relative to the mean baseline, weighted RMSE improved by 8.3% and weighted MAE by 7.6%. This means the demographic and social predictors are adding real predictive value, even if that improvement is not really that large.




The figure above compares the observed and predicted 'rwp_score' values on the internal test split. Each point represents a respondent, with their actual score on the x-axis and the model's predicted score on the y-axis. If the model predicted the outcome perfectly, the points would sit on a straight diagonal line where the predicted and observed values are equal. The figure shows that the model does move in the right direction. However, the predictions are much more tightly grouped around the middle of the scale than the actual scores. This is also shown by the predicted standard deviation ratio of 0.38, which means the spread of the predictions is much narrower than the spread of the actual scores. Overall, the model captures some real differences between respondents, but it is not very accurate at predicting individual scores. Overall, the results are modest. The model is clearly doing better than the baseline, but most of the variation in 'rwp_score' is still unexplained. This means the predictors are useful, but they are not enough to predict someone's score very accurately.

5.3 Trend Results

The final section looks at how the 'RWP_top_20' group changes over time and how it differs across social groups. This part is descriptive rather than predictive, so the aim is simply to show the main trends across the years.

5.3.1 Yearly Prevalence

Prevalence here means the share of respondents who fall into the 'RWP_top_20' group in each year. So a higher prevalence means that a larger proportion of respondents are above the top 20% cutoff, while a lower prevalence means that a smaller proportion are above it. This is not the same as the average 'rwp_score'. Instead, it shows how common the 'RWP_top_20' group is in each wave.




Prevalence of 'RWP_top_20' does not follow a simple linear trend. As you can see it starts at around 25% in 2011, drops a little in 2013 and then a lot to below 10% in 2019, and then is back up by 2024, this means that other than 2019 the amount of people exceeding the threshold is on average in the low to mid 20% range, this should not be read as a literal representation of how many right-wing populists exist in the wider population. What it shows is the share of respondents who score above the cutoff used for 'RWP_top_20', and could loosely be used as an indicator of the general levels of RWP or at least RWP leaning respondents.

The 2019 dip does not look like it is just down to poor data. The sample size and reliability results are still acceptable, with prevalence based on 2,595 respondents and Cronbach's alpha at 0.79 for the 'right' component and 0.79 for the 'welfare' component. The average scores across the different groups are also lower in 2019. This suggests that 2019 reflects a broader drop in the measure rather than simply just a data issue. I will come back to what might explain this in the Discussion chapter, where I will speak about the broader political context.

5.3.2 Main Predictors Over Time

The yearly regression summaries suggest that some predictors matter more consistently than others. As described in the Methods chapter, predictor strength in these yearly models is measured using grouped permutation importance, based on how much weighted R^2 drops when each predictor is shuffled.




Education is the clearest example. It is the strongest predictor in every year, which makes it the most stable result in the trend analysis. Religion and employment status are also consistently important, although they are less stable than education. Employment status is second strongest in 2013, while religion is second strongest in 2019 and 2024. Class also shows some variation, but less consistently, and it becomes noticeably weaker in 2024.


The yearly model-fit results are modest, so these predictors explain only part of the variation in ``rwp_score``. Even so, the results are useful because they show which social divisions are most consistently linked to the measure. Education is the strongest predictor in every year, while religion and employment status are also important but less consistent. This suggests some continuity over time, especially for education, but not one fixed social profile that is identical in every wave.

5.3.3 Subgroup Differences

The subgroup summaries make this divide easier to see directly. For education, the result is clear in every year. Respondents in the ‘low’ education group consistently score higher than those in the ‘high’ group, with the ‘mid’ group in between. This makes education the clearest and most stable result in the dissertation, suggesting that lower education is consistently linked to higher levels of right-wing populist attitudes.



Employment status also shows a clear divide. In every year, students have the lowest ‘rwp_score’, while retired respondents have the highest. This suggests that employment status is picking up an important social difference in the RWP measure. However, this should not be read only as an employment effect. It is also likely to reflect age and generational differences, since students are usually younger and retired respondents are usually older. So, while education remains the clearest result overall, employment status also shows a consistent difference across the years.



6 Discussion

6.1 Main Findings

Taken together, the results give a fairly clear picture. Demographic and social predictors do contain real information about 'rwp_score', but only a limited amount. The share of respondents in 'RWP_top_20' does not move in one straight direction across time, falling sharply in 2019 before rising again by 2024. Across the predictor results, education is the clearest and most consistent divide, while religion and employment status are also important but a little less stable.

These findings fit the literature reasonably well. The strong education result fits earlier work linking right-populist politics to educational and cultural divides, especially in the Brexit period [10, 11]. The limited predictive performance also fits the literature, because demographics are clearly only part of the broader story [15, 16].

6.2 What the Measure Captures

The first point to discuss is what this measure is really capturing. The final measure brings together distrust, authoritarian and immigration attitudes with welfare attitudes. That matters because it means the measure is not picking up one pure or perfectly

defined definition. Instead, it is picking up a bunch of related attitudes that overlap with the literature on RWP.

That is useful, but it is also a limitation. RWP is already difficult to define clearly in the literature [1, 3], and it is even harder to measure cleanly in survey data. Some of the attitudes used here overlap with ideas that are not uniquely RWP. Distrust of elites, for example, can also appear in left-wing populism. Welfare attitudes are also more complicated than the rest of the measure, because support for welfare can sit in different political traditions. So the final measure should be treated as a practical approximation built from the items the BSA makes available across the waves, not as a perfect measure of some fixed political belief.

This also matters for how the results are interpreted. The measure is strong enough to be useful, but it is still a constructed measure. It gives a workable way of tracing one right-populist attitude profile across time, but it cannot fully separate every type of anti-elite, nationalist, authoritarian, or welfare-based politics from one another. That fits the measurement literature, which repeatedly shows that populism is easier to discuss in theory than to measure neatly in practice [6, 8, 9].

6.3 Understanding the Main Social Differences

The clearest result in the dissertation is the education divide. Lower education is associated with higher `rwp_score` in every year, both in the yearly regression summaries and in the subgroup means. This fits closely with the literature reviewed earlier, especially work linking right-populist politics to cultural backlash, Brexit-era conflict, and divides between more and less educated groups [10, 11].

Employment status also shows a clear divide. Students have the lowest scores in every year, while retired respondents have the highest. This fits reasonably well with the literature reviewed earlier, which suggests age can matter here, although not always in a simple way [11, 23]. It also fits the current political context. Under Zack Polanski, the Greens have been described as taking a more left-populist or eco-populist direction [30], while Reform are a clear right-populist party in British politics [13]. Recent YouGov polling points to a strong age divide between them: Reform performs best among older voters, especially those in their 60s, while the Greens perform best among younger voters and are strongest among 18 to 24 year olds [31]. Age is not the same thing as employment status, so the comparison should be treated carefully. Even so, the overlap is useful here, since students are usually concentrated in younger age groups while retired respondents are usually concentrated in older age groups. This suggests that the employment-status result may show a wider age-linked and generational divide, which is reinforced by the current political context.

6.4 Understanding Change Over Time

One of the most interesting findings in the dissertation is how RWP changes across the years. The prevalence results do not show stability or a simple steady increase or decline. Instead, the measure falls sharply in 2019 and then rises again by 2024. This suggests that RWP in the BSA is not fixed across the period, but shifts with changing political conditions. The 2019 dip is especially important. As argued in the Results chapter, it does not look like a simple data problem. The sample and reliability results are still acceptable, and the average scores across the main groups are also lower in that year. So 2019 seems to reflect a broader drop in the measure rather than just weaker data.

One possible explanation is the political context of that year. Brexit was dominating British politics, but 2019 was also the year of a general election in which welfare and public spending were being brought into politics from more than one direction. There was still a visible anti-establishment and, in some respects, left-populist backing around Labour under Jeremy Corbyn [25], while Boris Johnson's campaign also placed heavy emphasis on NHS funding and spending promises [26, 27]. That matters because this measure includes welfare attitudes alongside distrust, authoritarian and immigration attitudes. In other words, welfare was politically important on both sides. It is therefore possible that some respondents who were anti-elite or anti-establishment in a broader sense were not captured as strongly by this right-populist measure because their welfare attitudes pulled them in a different direction. In that sense, the 2019 dip may reflect both a real change in the distribution of attitudes and a limit in what this particular measure can capture.

6.5 What the Model Performance Means

The modelling results help qualify the rest of the story. The ridge model clearly performs better than the mean baseline, so demographic and social characteristics do carry real information about where respondents sit on the score. That matters because it shows that RWP is not randomly distributed across the population.

At the same time, the model only explains a limited share of the overall variation. The predictions also remain too compressed toward the middle. So the model is useful, but only up to a point. It can pick up some real differences in who tends to score higher or lower, but it cannot predict the score very closely.

That also fits the literature. Research on populist attitudes regularly points to the importance of mistrust, efficacy, perceived deprivation, resentment, and broader political beliefs [15, 16]. This dissertation does not model all of that. It focuses on demographic and social predictors that can be harmonised across waves. So the results suggest that demographics matter, but they do not come close to explaining everything. That is not inherently a failure of the project. It shows that RWP is only partly

tied to background characteristics but also depends heavily on deeper political and attitudinal factors.

6.6 Main Limitations

The final RWP measure is shaped by the limits of the BSA. It is built from the survey items that are available across all the waves, rather than from a survey designed specifically to measure populism. Because of that, it is useful for this project, but it is still imperfect. Some of the attitudes included in the measure, especially anti-elite views and welfare attitudes, can also overlap with broader left-populist or anti-establishment politics. 'RWP_top_20' is also only a blunt summary of 'rwp_score'. It is useful for prevalence analysis, but it should not be treated as a direct measure of RWP. There are also limits to comparing the waves directly, because the BSA was not carried out in exactly the same way across the whole period, especially after the survey mode changed after 2020.

The main ridge model only explains a limited share of the variation in 'rwp_score', and the predicted scores are still pulled too much toward the middle. The yearly predictor rankings come from separate yearly models, so they are useful for showing which variables mattered most within each year, but they are not a final explanation of what causes RWP. Grouped permutation importance gives a useful way of comparing predictors across years, but it still depends on the fitted yearly models and the available predictor set.

A further limitation is that some social predictors, especially class and religion, may be harder to interpret cleanly over time because the measure does not always align neatly with the political mixes found in real groups. For example, some working-class respondents may be anti-elite, culturally conservative, or hostile to immigration while still being supportive of welfare, which means the welfare component can pull their score downward and understate some forms of working-class RWP. A similar issue applies to religion, where some respondents may hold socially conservative views linked to a traditional faith while also supporting welfare or other positions more commonly associated with the left. This makes class and religion informative, but also more difficult to interpret as simple or uniform predictors.

6.7 Future Work

A clear next step for future research would be to use survey data that is more directly tailored to right-wing populism. The main limitation of this dissertation is that the BSA does not contain a dedicated right-wing populism scale repeated consistently across many years, so the final measure had to be built as an approximation from what was available. Future work could therefore apply the same general approach to data with more targeted questions, which would allow a more direct and potentially more accurate measure of RWP.

This could be achieved through a purpose-built survey, or through work with organisations such as NatCen to include more politically focused questions in future survey waves. That would also help with the predictor set, which in this dissertation was limited by missing data. A better and wider-reaching survey could therefore improve both the measurement of RWP and the prediction stage of the analysis.

A final area for future work would be to test different models. This dissertation used ridge regression, but future work could try out more complex models to see whether they can predict RWP more effectively. More complex models may be better at finding relationships that a linear model cannot. Therefore, trying them out would help show whether the modest results in this dissertation come from the predictor set itself, or from the choice of model. Ultimately, improving both would hopefully lead to better results in the future.

6.8 Final Remarks

Even with those limits, the dissertation still makes a useful contribution to the topic. It builds a repeatable survey-based measure of RWP from repeated BSA data, shows how that measure changes across the years, and identifies which social divides are most consistent across waves. The clearest of those is education, while the results over time show that RWP in Britain does not move in one simple direction.

More broadly, the dissertation connects the literature to a practical empirical design. It does not claim to produce a perfect measure of RWP, or to explain it fully through demographics alone. What it does show is that RWP can be partially studied in a transparent and repeatable way using the BSA data, and that it remains politically relevant in a Britain still shaped by Brexit and wider political fragmentation.

Appendix

Bibliography

[1] Mudde C. The populist zeitgeist. *Government and Opposition*. 2004;39(4):541–563. doi:10.1111/j.1477-7053.2004.00135.x.

[2] Moffitt B. *The Global Rise of Populism: Performance, Political Style, and Representation*. Stanford, CA: Stanford University Press; 2016.

[3] Mudde C. *Populist Radical Right Parties in Europe*. Cambridge: Cambridge University Press; 2007.

[4] Akkerman A, Mudde C, Zaslove A. How populist are the people? Measuring populist attitudes in voters. *Comparative Political Studies*. 2014;47(9):1324–1353. doi:10.1177/0010414013512600.

- [5] Schulz A, Müller P, Schemer C, Wirz DS, Wettstein M, Wirth W. Measuring populist attitudes on three dimensions. *International Journal of Public Opinion Research*. 2018;30(2):316–326. doi:10.1093/ijpor/edw037.
- [6] Wuttke A, Schimpf CH, Schoen H. When the whole is greater than the sum of its parts: On the conceptualization and measurement of populist attitudes and other multidimensional constructs. *American Political Science Review*. 2020;114(2):356–374. doi:10.1017/S0003055419000807.
- [7] Geurkink B, Zaslove A, Sluiter R, Jacobs K. Populist attitudes, political trust, and external political efficacy: Old wine in new bottles? *Political Studies*. 2020;68(1):247–267. doi:10.1177/0032321719842768.
- [8] Wettstein M, Schulz A, Steenbergen M, Schemer C, Müller P, Wirz DS, Wirth W. Measuring populism across nations: Testing for measurement invariance of an inventory of populist attitudes. *International Journal of Public Opinion Research*. 2020;32(2):284–305. doi:10.1093/ijpor/edz018.
- [9] Jungkunz S, Fahey RA, Hino A. How populist attitudes scales fail to capture support for populists in power. *PLoS ONE*. 2021;16(12):e0261658. doi:10.1371/journal.pone.0261658.
- [10] Norris P, Inglehart R. *Cultural Backlash: Trump, Brexit, and Authoritarian Populism*. Cambridge: Cambridge University Press; 2019.
- [11] Hobolt SB. The Brexit vote: a divided nation, a divided continent. *Journal of European Public Policy*. 2016;23(9):1259–1277. doi:10.1080/13501763.2016.1225785.
- [12] Curtice J, Montagu I, Sivathasan C. Damaged politics? The impact of the 2019–24 Parliament on political trust and confidence. In: Clery E, Curtice J, Jessop C, eds. *British Social Attitudes: The 41st Report*. London: National Centre for Social Research; 2024.
- [13] Hayton, R. To conserve or to reform? The reshaping of the right in British politics. *Br Polit* 20, 381–399 (2025). <https://doi.org/10.1057/s41293-025-00287-4>
- [14] Schimpf CH, Wuttke A, Schoen H. Neither a trait nor wildly fluctuating: On the stability of populist attitudes and its implications for empirical research. *British Journal of Political Science*. 2024;54(3):979–992. doi:10.1017/S0007123423000492.
- [15] Piterová I, Loziak A. A comprehensive model for predicting populist attitudes. *Journal of Social and Political Psychology*. 2024;12(1):73–88. doi:10.5964/jspp.11539.
- [16] Loziak A, Piterová I. The relationship between populist attitudes, political efficacy, trust and cognitive beliefs: A network analysis approach. *Slovak Journal of Political Sciences*. 2024;24(1):6–30. doi:10.34135/sjps.240101.

- [17] Cordero G, Zagórski P, Rama J. Give me your least educated: Immigration, education and support for populist radical right parties in Europe. *Political Studies Review*. 2022;20(3):517–524. doi:10.1177/14789299211029110.
- [18] Ciccolini G. Who is left behind? Economic status loss and populist radical right voting. *Perspectives on Politics*. 2025;First View:1–18. doi:10.1017/S1537592724002858.
- [19] Coffé H. Gender, gendered personality traits and radical right populist voting. *Politics*. 2019;39(2):170–185. doi:10.1177/0263395717745476.
- [20] Marcinkiewicz K, Dassonneville R. Do religious voters support populist radical right parties? Opposite effects in Western and East-Central Europe. *Party Politics*. 2022;28(3):444–456. doi:10.1177/1354068820985187.
- [21] Rickardsson J. The urban–rural divide in radical right populist support: the role of resident’s characteristics, urbanization trends and public service supply. *The Annals of Regional Science*. 2021;67:211–242. doi:10.1007/s00168-021-01046-1.
- [22] Haugsgjerd A, Bergh J. Explaining the Surge of the Populist Radical Right: A Time-Series Analysis of the Effects of Immigration and the Economy in Norway. *Political Behavior*. 2024;46:1633–1652. doi:10.1007/s11109-023-09887-6.
- [23] Almstedt Valldor A. Radical right support and the deepened rural-urban and educational divides among younger generations. *Acta Sociologica*. Published online June 25, 2025. doi:10.1177/00016993251348067.
- [24] Ladreit C. Austerity and young people’s political attitudes in the UK. *Journal of Public Policy*. 2023;43(4):603–636. doi:10.1017/S0143814X23000107.
- [25] Watts J, Bale T. Populism as an intra-party phenomenon: The British Labour Party under Jeremy Corbyn. *The British Journal of Politics and International Relations*. 2019;21(1):99–115. doi:10.1177/1369148118806115.
- [26] Labour Party. It’s Time for Real Change: The Labour Party Manifesto 2019. London: Labour Party; 2019. <https://www.labourinternational.net/wp-content/uploads/sites/70/2020/09/Real-Change-Labour-Manifesto-2019.pdf>
- [27] Conservative and Unionist Party. Get Brexit Done: Unleash Britain’s Potential. London: Conservative and Unionist Party; 2019. <https://cchq2019.webflow.io/our-plan>
- [28] YouGov. Voting intention, 29–30 March 2026: Ref 23%, Grn 19%, Con 19%, Lab 18%, LD 12%. London: YouGov; 2026. <https://yougov.com/en-gb/articles/54444-voting-intention-29-30-march-2026-ref-23-grn-19-con-19-lab-18-ld-12>
- [29] Simpson I. Polling breakdown from March 2026: Latest polls see continued fragmentation. London: Electoral Reform Society; 2026. <https://electoral->

reform.org.uk/polling-breakdown-from-march-2026-latest-polls-see-continued-fragmentation/

[30] The Guardian. 'I thought politics was a dirty thing': Zack Polanski on his 'eco-populist' vision for the Green party. The Guardian. 13 May 2025. <https://www.theguardian.com/politics/2025/may/13/i-thought-politics-was-a-dirty-thing-zack-polanski-on-his-eco-populist-vision-for-the-green-party>

[31] Difford D. How would Britain vote at the start of 2026? London: YouGov; 2026. <https://yougov.com/en-gb/articles/53923-how-would-britain-vote-at-the-start-of-2026>